

How this wallet system can scale in production

1. **Business logic is isolated in the service**
 - All wallet rules (funding, transfers) live in WalletsService.
 - Controllers are thin, so scaling the API layer doesn't affect business logic.
 - This makes it easy to add more endpoints or versions without breaking core behavior.
 2. **Database transactions guarantee consistency**
 - Transfers run inside a database transaction.
 - Either both wallets are updated, or nothing is.
 - This prevents partial updates and keeps balances consistent even under load.
 3. **Idempotency protects against retries**
 - Each fund or transfer request uses an idempotency key.
 - If a client retries the same request, the system returns the previous result instead of processing again.
 - This is critical for production systems where network retries are common.
 4. **Balances are stored as decimals, not floats**

Wallet balances are stored as strings backed by database decimals.

 - This avoids rounding errors and makes the system safe for financial calculations at scale.
 5. **Writes are controlled, reads can scale independently**
 - Wallet updates are strongly consistent and transactional.
 - Read operations (like wallet history) can later be moved to read replicas without changing the code.
 - This allows horizontal scaling as traffic grows.
 6. **Wallet history enables auditing and reconciliation**
 - Every balance change creates a history record.
 - This makes it easy to debug issues, reconcile accounts, or run analytics.
 - In production, this is essential for trust and compliance.
1. **Stateless API design**
 - The service does not keep in-memory state.
 - Any instance can handle any request.
 - This makes horizontal scaling straightforward using load balancers.
 2. **Easy to extend without redesign**
 - New features like limits, fees, reversals, or multi-currency wallets can be added without changing the core flow.

- The existing transaction and history structure already supports this.
3. **Clear failure handling**
 - Errors like insufficient balance or invalid wallets are caught early.
 - Failed operations do not leave dirty data behind.
 - This keeps the system stable under high traffic.
 4. **Path to production hardening**
 - Row locking can be added to prevent race conditions.
 - Caching can be introduced for performance.
 - The current design does not block any of these improvements.